

Locality Sensitive Hashing

Clément King

15 août 2017

Summary

Problem and Motivations

Locality Sensitive Hashing

Algorithm

Random Binary Projections

Code

Problem and Motivations (1)


Nearest Neighbor (NN) problem : given a collection of n points, build a data structure which, given a query point, reports the point that is the closest to the queried one.

In general points are represented as vectors in \mathbb{R}^d .

Issue : Naive solutions suffer from the "curse of dimensionality". Brute force linear search is too long for big database. To compute all the distances between n points, one needs to compute $\frac{n(n-1)}{2}$ distances.

Problem and Motivation (2)

y-axis : computation time, x-axis : number of points



Alternative to brute force ? Near Neighbor search !

Summary

Problem and Motivations

Locality Sensitive Hashing

Algorithm

Random Binary Projections

Code

Locality Sensitive Hashing

The algorithm relies on the existence of hash functions (function that map data of arbitrary size to a fixed size subspace).

Definition : A family \mathbb{H} is said (R, cR, p_1, p_2) - sensitive if for any points p and q in \mathbb{R}^d , and any hash function h in \mathbb{H} :

- $P_{h \in \mathbb{H}}(h(p) = h(q)) \geq p_1$ if $d(p, q) \leq R$
- $P_{h \in \mathbb{H}}(h(p) = h(q)) \leq p_2$ if $d(p, q) \geq cR$

where $c > 1$, $p_1 > p_2$, R is some radius and d is some distance in \mathbb{R}^d .

What does that mean ? We want the probability of collision (p_1) to be higher for objects that are close to each other ($d(p, q) \leq R$) than for those that are far apart ($d(p, q) \geq cR$).

Example

Suppose our data points are binary (10001, 11011, ...), and let d be the Hamming distance (i.e $d(101, 100) = 1$, the number of bits that are different).

Now, if $\mathbb{H} = \{h_i : \{0, 1\}^d \rightarrow \{0, 1\} | h_i(p) = p_i\}$, choosing a random h in \mathbb{H} means that $h(p)$ returns a random coordinate of $p \in \mathbb{R}^d$.

Is that a "sensitive" family? We have that $P_{h \in \mathbb{H}}(h(p) = h(q))$ is equal to the number of coordinates on which p and q agree divided by the total number of coordinates.

Hence $p_1 = 1 - \text{proportion of bits that differ} = 1 - \frac{R}{d}$ and similarly, $p_2 = 1 - \frac{cR}{d}$.

\Rightarrow As long as $c > 1$, $p_1 > p_2$.

Summary

Problem and Motivations

Locality Sensitive Hashing

Algorithm

Random Binary Projections

Code

Amplification (1)

Since p_1 and p_2 may be close to each other, we cannot use a hash function by itself.

Consider k and $B \in \mathbb{N}$, and define :

$$g_i(q) = (h_{1,i}(q), \dots, h_{k,i}(q)) \text{ for all } 1 \leq i \leq B,$$

where each $h_{j,i}$ is chosen independently in \mathbb{H} .

$g_i : \mathbb{R}^d \rightarrow \mathbb{Z}^k$ (we created a hash function using k random hash functions) and so each query point q is placed into B different buckets.

To process a query point q we look at the B buckets q is in and compute the distances between the points of those buckets and q .

Amplification (2)

Two strategies :

- Stopping the search in the buckets after scanning at most $3B$ collisions (see Indyk-Motwani).
- Scan all the collisions (that's what we do) : Let p be any point in the R - neighborhood of q .

For any g_i , $P(g_i(p) = g_i(q)) \geq p_1^k$ and so,

$P(g_i(p) = g_i(q) \text{ for some } i)$

$= P(\exists i \text{ such that } g_i(p) = g_i(q))$

$= 1 - P(\nexists i \text{ such that } g_i(p) = g_i(q))$

$= 1 - P(\{g_i(p) \neq g_i(q)\} \text{ for all } i)$


$\geq 1 - (1 - p_1^k)^B$

Algorithm

- Init
 - Choose B functions $g_i : \mathbb{R}^d \rightarrow \mathbb{Z}^k$
 - Construct B hash tables where the i^{th} table contains the points hashed with the $g_i^t h$ function
- Algo
 - Input : a query point q
 - For $i = 1, \dots, B$
 - Retrieve points from the bucket $g_i(q)$ in the i^{th} hash table
 - Compute the distances between all the points and q

Trade-off

Large values of k lead to larger gap between p_1 and p_2 which causes the hashing to be selective (good!). However, if k is large p_1^k is small (very bad!!), and so we need a big B to compensate.



Left : proba that $g_i(p) = g_i(q)$ for some i for different values of k (the larger k is, the more we partition our space)

Right : proba that $g_i(p) = g_i(q)$ for any i for different values of k (the larger B is, the more buckets we're scanning).

Summary

Problem and Motivations

Locality Sensitive Hashing

Algorithm

Random Binary Projections

Code

Random Binary Projections (1)


Let our space be the unit sphere in \mathbb{R}^d , say \mathbb{S} .

- For i from 1 to k
 - Choose a unit vector u_i at random in \mathbb{S}
 - For each point v we have, $g_i(v) = \text{sgn}(u_i \cdot v)$.

We get a signature matrix M of size $k \times (\text{number of points } v)$ such that $M(i, v) = \text{sgn}(u_i \cdot v)_{i \leq k}$.

Hence, the candidates to be nearest neighbors of a given input vector w are the vectors for which the column of the matrix are identical to the hash representation of w .

Random Binary Projection (2)



It happens that $P(\text{sgn}(u \cdot v) = \text{sgn}(u \cdot w)) = 1 - \frac{\arccos(v \cdot w)}{\pi}$

Summary

Problem and Motivations

Locality Sensitive Hashing

Algorithm

Random Binary Projections

Code

Code(1)

```
# https://github.com/pixelogik/NearPy
from nearpy import Engine
from nearpy.distances import CosineDistance
from nearpy.hashes import RandomBinaryProjections

# Let matrix be the scoring vector matrix, i.e each line
# corresponds to a score vector

lsh = RandomBinaryProjections("rbp0", nb_random_vectors)
engine = Engine(matrix.shape[1],
                lshashes=[lsh], distance=CosineDistance())
```

Code (2)

```
mapping = defaultdict(list)
for i in xrange(matrix.shape[0]):
    scoring_vector = matrix[i, :]
    hash = lsh.hash_vector(scoring_vector[0])
    mapping[hash_v].append(i)
```

```
# output is like:
# {1001: [5785, 485795, 585, ...],
#  0010: [ 27402, 1957, ...],
#  ... }
```

Andoni, A., Indyk, P., *Near-Optimal Hashing Algorithms for Approximate Nearest Neighbor in High Dimensions*

Indyk, P., Motwani, R., ANN : Towards removing the curse of dimensionality